

P1: Hive

William Terry

August 2021

Problem1

Problem Scenario 1:

What is the total number of consumers for Branch1?

Answer: 1115974

What is the number of consumers for the Branch2?

Answer: 5099141

```

create table BevA(Beverage STRING, BranchID String)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_BranchA.txt' OVERWRITE INTO TABLE BevA;

create table BevB(Beverage STRING, BranchID String)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_BranchB.txt' OVERWRITE INTO TABLE BevB;

create table BevC(Beverage STRING, BranchID String)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_BranchC.txt' OVERWRITE INTO TABLE BevC;

create table ConsA(Beverage STRING, Consumed INT)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_ConscountA.txt' OVERWRITE INTO TABLE ConsA;

create table ConsB(Beverage STRING, Consumed INT)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_ConscountB.txt' OVERWRITE INTO TABLE ConsB;

create table ConsC(Beverage STRING, Consumed INT)row format delimited fields terminated by ',' stored as textfile;
LOAD DATA INPATH '/user/williamgantterry/Bev_ConscountC.txt' OVERWRITE INTO TABLE ConsC;

CREATE TABLE BevAll(Beverage STRING,BranchID STRING)row format delimited fields terminated by ',' stored as textfile;
INSERT INTO TABLE BevAll (SELECT * FROM BevA UNION ALL SELECT * FROM BevB UNION ALL SELECT * FROM BevC);

CREATE TABLE ConsAll(Beverage STRING,Consumed INT) row format delimited fields terminated by ',' stored as textfile;
INSERT INTO TABLE ConsAll(SELECT * FROM ConsA UNION ALL SELECT * FROM ConsB UNION ALL SELECT * FROM ConsC);

CREATE TABLE TotalCons(BranchID STRING, TotalConsumers INT);
INSERT INTO TABLE TotalCons(
SELECT BranchID, sum(ConsAll.Consumed) FROM BevAll join ConsAll On (ConsAll.Beverage=BevAll.Beverage) GROUP BY BevAll.BranchID);

SELECT * FROM TotalCons;

```

totalcons.branchid	totalcons.totalconsumers
Branch1	1115974
Branch2	5099141
Branch3	2591062
Branch4	3940862
Branch5	3941177
Branch6	8032799
Branch7	6566187
Branch8	2583583
Branch9	5208016

williamgantterry@LAPTOP-H2L556UK: ~

9 rows selected (0.122 seconds)

0: jdbc:hive2://> SELECT * FROM TotalCons;

OK

totalcons.branchid	totalcons.totalconsumers
Branch1	1115974
Branch2	5099141
Branch3	2591062
Branch4	3940862
Branch5	3941177
Branch6	8032799
Branch7	6566187
Branch8	2583583
Branch9	5208016

9 rows selected (0.107 seconds)

0: jdbc:hive2://> █

Problem2

Problem Scenario 2:

What is the most consumed beverage on Branch1?

Answer: Special Cappucino with 108163

What is the least consumed beverage on Branch2? Answer: Cold Mocha with 47524

What is the Average (median) consumed beverage of Branch2

Answer: 26/51 in order which is Small Latte with 93184

```
SELECT BranchID, sum(ConsAll.Consumed),BevAll.Beverage FROM BevAll join ConsAll On (ConsAll.Beverage=BevAll.Beverage)
WHERE BranchID="Branch1" GROUP BY BevAll.BranchID,BevAll.Beverage ORDER BY sum(ConsAll.Consumed) DESC]
```

```
;
```

branchid	_c1	bevall.beverage
Branch1	108163	Special_cappuccino
Branch1	103894	LARGE_cappuccino
Branch1	54378	Double_Espresso
Branch1	54208	Special_Coffee
Branch1	53824	LARGE_Espresso
Branch1	52005	Cold_Lite
Branch1	51604	MED_MOCHA
Branch1	50820	ICY_Espresso
Branch1	50644	Triple_MOCHA
Branch1	50420	Mild_LATTE
Branch1	50228	Double_MOCHA
Branch1	50202	SMALL_Espresso
Branch1	50021	Triple_LATTE
Branch1	49482	ICY_MOCHA
Branch1	49237	SMALL_MOCHA
Branch1	48283	Special_Espresso
Branch1	47924	Cold_LATTE
Branch1	47211	MED_LATTE
Branch1	46834	Mild_Lite
Branch1	46592	SMALL_LATTE

Figure: Problem 2

williamgantterry@LAPTOP-H2L556UK: ~

Ended Job = job_1627998776566_0042

MapReduce Jobs Launched:

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 5.98 sec HDFS Read: 123682 HDFS Write: 760 SUCCESS

Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 3.73 sec HDFS Read: 8450 HDFS Write: 871 SUCCESS

Total MapReduce CPU Time Spent: 9 seconds 710 msec

OK

branchid	_c1	bevall.beverage
Branch1	108163	Special_cappuccino
Branch1	103894	LARGE_cappuccino
Branch1	54378	Double_Espresso
Branch1	54208	Special_Coffee
Branch1	53824	LARGE_Espresso
Branch1	52005	Cold_Lite
Branch1	51604	MED_MOCHA
Branch1	50820	ICY_Espresso
Branch1	50644	Triple_MOCHA
Branch1	50420	Mild_LATTE
Branch1	50228	Double_MOCHA
Branch1	50202	SMALL_Espresso
Branch1	50021	Triple_LATTE
Branch1	49482	ICY_MOCHA
Branch1	49237	SMALL_MOCHA
Branch1	48283	Special_Espresso
Branch1	47924	Cold_LATTE
Branch1	47211	MED_LATTE
Branch1	46834	Mild_Lite
Branch1	46592	SMALL_LATTE

20 rows selected (61.25 seconds)

0: jdbc:hive2://>

```
SELECT BranchID, ROW_NUMBER() OVER (ORDER BY sum(ConsAll.Consumed)) AS row_num, sum(ConsAll.Consumed), BevAll.Beverage
FROM BevAll join ConsAll On (ConsAll.Beverage=BevAll.Beverage) WHERE BranchID="Branch2"
GROUP BY BevAll.BranchID, BevAll.Beverage ORDER BY sum(ConsAll.Consumed) DESC
;
```

branchid	row_num	_c2	bevall.beverage
Branch2	51	328074	Mild_cappuccino
Branch2	50	285309	MED_cappuccino
Branch2	49	284943	Cold_cappuccino
Branch2	48	216326	Special_cappuccino
Branch2	47	207788	LARGE_cappuccino
Branch2	46	204476	ICY_cappuccino
Branch2	45	202854	SMALL_cappuccino
Branch2	44	197620	Double_cappuccino
Branch2	43	193636	Triple_cappuccino
Branch2	42	112378	Double_LATTE
Branch2	41	108292	Special_MOCHA
Branch2	40	107648	LARGE_Espresso
Branch2	39	104466	LARGE_LATTE
Branch2	38	104010	Cold_Lite
Branch2	37	103572	MED_Lite
Branch2	36	101192	LARGE_MOCHA
Branch2	35	100408	Triple_Lite
Branch2	34	100042	Triple_LATTE
Branch2	33	98964	ICY_MOCHA
Branch2	32	98614	SMALL_Lite
Branch2	31	98532	Double_Lite
Branch2	30	98332	ICY_LATTE
Branch2	29	96566	Special_Espresso
Branch2	28	94422	MED_LATTE
Branch2	27	93668	Mild_Lite
Branch2	26	93184	SMALL_LATTE
Branch2	25	54378	Double_Espresso
Branch2	24	54208	Special_Coffee
Branch2	23	54050	Special_Lite
Branch2	22	53681	Mild_MOCHA
Branch2	21	51932	Cold_Coffee
Branch2	20	51880	Mild_Coffee
Branch2	19	51694	Double_Coffee
Branch2	18	51647	LARGE_Coffee
Branch2	17	51604	MED_MOCHA

-	Branch2	16	50820	ICY_Espresso	
-	Branch2	15	50648	Triple_Coffee	
-	Branch2	14	50644	Triple_MOCHA	
-	Branch2	13	50604	Special_LATTE	
-	Branch2	12	50228	Double_MOCHA	
-	Branch2	11	50202	SMALL_Espresso	
-	Branch2	10	49501	SMALL_Coffee	
-	Branch2	9	49237	SMALL_MOCHA	
-	Branch2	8	48994	MED_Coffee	
-	Branch2	7	48775	Triple_Espresso	
-	Branch2	6	48555	ICY_Coffee	
-	Branch2	5	48462	Mild_Espresso	
-	Branch2	4	48413	MED_Espresso	
-	Branch2	3	48220	ICY_Lite	
-	Branch2	2	47924	Cold_LATTE	
-	Branch2	1	47524	Cold_MOCHA	
+-----+-----+-----+-----+					

Figure: Problem 2

```

williamganterry@LAPTOP-H2L556UK: ~
Total MapReduce CPU Time Spent: 15 seconds 170 msec
OK

```

branchid	row_num	_c2	bevall.beverage
Branch2	51	328074	Mild_cappuccino
Branch2	50	285309	MED_cappuccino
Branch2	49	284943	Cold_cappuccino
Branch2	48	216326	Special_cappuccino
Branch2	47	207788	LARGE_cappuccino
Branch2	46	204476	ICY_cappuccino
Branch2	45	202854	SMALL_cappuccino
Branch2	44	197620	Double_cappuccino
Branch2	43	193636	Triple_cappuccino
Branch2	42	112378	Double_LATTE
Branch2	41	108292	Special_MOCHA
Branch2	40	107648	LARGE_Espresso
Branch2	39	104466	LARGE_LATTE
Branch2	38	104010	Cold_Lite
Branch2	37	103572	MED_Lite
Branch2	36	101192	LARGE_MOCHA
Branch2	35	100408	Triple_Lite
Branch2	34	100042	Triple_LATTE
Branch2	33	98964	ICY_MOCHA
Branch2	32	98614	SMALL_Lite
Branch2	31	98532	Double_Lite
Branch2	30	98332	ICY_LATTE
Branch2	29	96566	Special_Espresso
Branch2	28	94422	MED_LATTE
Branch2	27	93668	Mild_Lite
Branch2	26	93184	SMALL_LATTE
Branch2	25	54378	Double_Espresso
Branch2	24	54208	Special_Coffee

Figure: Problem 2

Branch2	23	54950	Special_Lite
Branch2	22	53681	Mild_MOCHA
Branch2	21	51932	Cold_Coffee
Branch2	20	51880	Mild_Coffee
Branch2	19	51694	Double_Coffee
Branch2	18	51647	LARGE_Coffee
Branch2	17	51604	MED_MOCHA
Branch2	16	50820	ICY_Espresso
Branch2	15	50648	Triple_Coffee
Branch2	14	50644	Triple_MOCHA
Branch2	13	50604	Special_LATTE
Branch2	12	50228	Double_MOCHA
Branch2	11	50202	SMALL_Espresso
Branch2	10	49501	SMALL_Coffee
Branch2	9	49237	SMALL_MOCHA
Branch2	8	48994	MED_Coffee
Branch2	7	48775	Triple_Espresso
Branch2	6	48555	ICY_Coffee
Branch2	5	48462	Mild_Espresso
Branch2	4	48413	MED_Espresso
Branch2	3	48220	ICY_Lite
Branch2	2	47924	Cold_LATTE
Branch2	1	47524	Cold_MOCHA

-----+-----+-----+-----+
 51 rows selected (102.85 seconds)

Figure: Problem 2

Problem3

Problem Scenario 3 What are the beverages available on Branch10, Branch8, and Branch1? What are the common beverages available in Branch4,Branch7?

```
SELECT DISTINCT BevAll.beverage FROM BevAll WHERE BRANCHID="Branch1" or BRANCHID="Branch8" or BRANCHID="Branch10";
```

```
SELECT BevAll.beverage, BevAll.BRANCHID FROM BevAll WHERE BRANCHID="Branch1" or BRANCHID="Branch8" or BRANCHID="Branch10"
ORDER BY BevAll.beverage;
```

```
+-----+
| bevall.beverage |
+-----+
| Cold_Coffee     |
| Cold_LATTE      |
| Cold_Lite       |
| Cold_cappuccino |
| Double_Coffee   |
| Double_Espresso |
| Double_LATTE    |
| Double_MOCHA    |
| Double_cappuccino |
| ICY_Coffee      |
| ICY_Espresso    |
| ICY_Lite        |
| ICY_MOCHA       |
| ICY_cappuccino  |
| LARGE_Coffee    |
| LARGE_Espresso  |
| LARGE_MOCHA     |
| LARGE_cappuccino |
| MED_Coffee      |
| MED_Espresso    |
| MED_LATTE       |
| MED_MOCHA       |
| MED_cappuccino  |
| Mild_Coffee     |
| Mild_Espresso   |
| Mild_LATTE      |
| Mild_Lite       |
| Mild_cappuccino |
| SMALL_Espresso  |
| SMALL_LATTE     |
| SMALL_Lite      |
| SMALL_MOCHA     |
```

```
| Special_Espresso |  
| Special_LATTE  |  
| Special_Lite    |  
| Special_MOCHA   |  
| Special_cappuccino |  
| Triple_Coffee   |  
| Triple_Espresso  |  
| Triple_LATTE    |  
| Triple_Lite     |  
| Triple_MOCHA    |  
| Triple_cappuccino |  
+-----+-----+
```

Figure: Problem 3

--What are the common beverages available in Branch4,Branch7?

```
select BevAll.beverage from BevAll where BevAll.BRANCHID = "Branch4" and BevAll.beverage in (select BevAll.beverage from BevAll where BevAll.BRANCHID = "Branch7")
GROUP BY BevAll.beverage ;
```

bevall.beverage
Cold_Coffee
Cold_Espresso
Cold_LATTE
Cold_lite
Cold_MOCHA
Cold_cappuccino
Double_Coffee
Double_Espresso
Double_lite
Double_MOCHA
Double_cappuccino
ICY_Coffee
ICY_LATTE
ICY_lite
ICY_MOCHA
ICY_cappuccino
LARGE_Coffee
LARGE_Espresso
LARGE_LATTE
LARGE_lite
LARGE_MOCHA
LARGE_cappuccino
MED_Coffee
MED_Espresso
MED_lite
MED_MOCHA
MED_cappuccino
Mild_Coffee
Mild_Espresso
Mild_LATTE
Mild_lite
Mild_MOCHA
Mild_cappuccino

Figure: Problem 3

Problem1
Problem2
Problem3
Problem4
Problem5
Problem6
GitHub

```
- | Mild_Coffee |  
- | Mild_Espresso |  
- | Mild_LATTE |  
- | Mild_Lite |  
- | Mild_MOCHA |  
- | Mild_cappuccino |  
- | SMALL_Coffee |  
- | SMALL_Espresso |  
- | SMALL_LATTE |  
- | SMALL_Lite |  
- | SMALL_MOCHA |  
- | SMALL_cappuccino |  
- | Special_Coffee |  
- | Special_Espresso |  
- | Special_LATTE |  
- | Special_Lite |  
- | Special_MOCHA |  
- | Special_cappuccino |  
- | Triple_Coffee |  
- | Triple_Espresso |  
- | Triple_LATTE |  
- | Triple_Lite |  
- | Triple_MOCHA |  
- | Triple_cappuccino |  
+-----+
```

Figure: Problem 3

 williamgantterry@LAPTOP-H2L556UK: ~

```
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.67 sec HDFS Read: 39661 HDFS Write: 1203 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 670 msec
OK
```

```
+-----+
| bevall.beverage |
+-----+
| Cold_Coffee      |
| Cold_LATTE       |
| Cold_Lite        |
| Cold_cappuccino  |
| Double_Coffee    |
| Double_Espresso  |
| Double_LATTE     |
| Double_MOCHA     |
| Double_cappuccino|
| ICY_Coffee       |
| ICY_Espresso     |
| ICY_Lite         |
| ICY_MOCHA        |
| ICY_cappuccino   |
| LARGE_Coffee     |
| LARGE_Espresso   |
| LARGE_MOCHA      |
| LARGE_cappuccino |
| MED_Coffee       |
| MED_Espresso     |
| MED_LATTE        |
| MED_MOCHA        |
| MED_cappuccino   |
| Mild_Coffee      |
| Mild_Espresso    |
| Mild_LATTE       |
| Mild_Lite        |
+-----+
```

Problem4

```
-Problem Scenario 4
~create a partition, View for the scenario3.

Create VIEW SampView AS select BevAll.beverage from BevAll where BevAll.BRANCHID = "Branch4" and BevAll.beverage in
(select BevAll.beverage from BevAll where BevAll.BRANCHID = "Branch7")
GROUP BY BevAll.beverage;

SELECT * FROM SampView;

create table Part4(beverage STRING) PARTITIONED by (branchid STRING) row format delimited fields terminated by ',' stored as textfile;
INSERT INTO TABLE Part4(SELECT BevAll.beverage, BevAll.BRANCHID FROM BevAll WHERE BRANCHID="Branch1" or BRANCHID="Branch8" or BRANCHID="Branch10"
ORDER BY BevAll.beverage);

show partitions Part4;

+-----+
-|      partition      |
+-----+
-| branchid=Branch1    |
-| branchid=Branch8    |
+-----+
```

Figure: Problem 4

Problem4

```
williamgantterry@LAPTOP-H2L556UK: ~  
Ended Job = job_1627998776566_0030  
MapReduce Jobs Launched:  
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 6.01 sec HDFS Read: 43448 HDFS Write: 1371 SUCCESS  
Total MapReduce CPU Time Spent: 6 seconds 10 msec  
OK  
+-----+  
| sampview.beverage |  
+-----+  
| Cold_Coffee        |  
| Cold_Espresso      |  
| Cold_LATTE         |  
| Cold_Lite          |  
| Cold_MOCHA         |  
| Cold_cappuccino     |  
| Double_Coffee      |  
| Double_Espresso     |  
| Double_Lite        |  
| Double_MOCHA       |  
| Double_cappuccino   |  
| ICY_Coffee         |  
| ICY_LATTE          |  
| ICY_Lite           |  
| ICY_MOCHA          |  
| ICY_cappuccino     |  
| LARGE_Coffee       |  
| LARGE_Espresso     |  
| LARGE_LATTE        |  
| LARGE_Lite         |
```

Problem4



williamgantterry@LAPTOP-H2L556UK: ~

```
51 rows selected (34.305 seconds)
0: jdbc:hive2://> show partitions Part4;
OK
+-----+
| partition |
+-----+
| branchid=Branch1 |
| branchid=Branch8 |
+-----+
2 rows selected (0.096 seconds)
```

Problem5

-Problem Scenario 5

-Alter the table properties to add "note","comment"

```
ALTER TABLE Part4 SET TBLPROPERTIES ('note'='this is a note');
```

```
ALTER TABLE Part4 SET TBLPROPERTIES ('comment'='this is a comment');
```

-Show tblproperties Part4;

prpt_name	prpt_value
bucketing_version	2
comment	this is a comment
last_modified_by	williamganterterry
last_modified_time	1628007015
note	this is a note
transient_lastDdlTime	1628007015

Figure: Problem 5

Problem5

```
williamgantterry@LAPTOP-H2L556UK: ~  
8 rows selected (0.223 seconds)  
0: jdbc:hive2://> Show tblproperties Part4;  
OK  
+-----+-----+  
|      prpt_name      |      prpt_value      |  
+-----+-----+  
| bucketing_version   | 2                     |  
| comment              | this is a comment    |  
| last_modified_by     | williamgantterry     |  
| last_modified_time   | 1628007015            |  
| note                 | this is a note       |  
| transient_lastDdlTime | 1628007015            |  
+-----+-----+  
6 rows selected (0.09 seconds)  
0: jdbc:hive2://>
```

Problem6

Problem Scenario 6

-Remove a row from the any Scenario.

```
DROP TABLE IF EXISTS Tab62;
```

```
CREATE TABLE Tab62(BranchID STRING, TotalConsumers INT) STORED AS ORC;  
INSERT INTO TABLE Tab62(SELECT * FROM TotalCons);
```

```
SELECT * FROM Tab62;
```

+-----+-----+	
- tab62.branchid	tab62.totalconsumers
+-----+-----+	
- Branch1	1115974
- Branch2	5099141
- Branch3	2591062
- Branch4	3940862
- Branch5	3941177
- Branch6	8032799
- Branch7	6566187
- Branch8	2583583
- Branch9	5208016
+-----+-----+	

Figure: Problem 6

Problem6

```
SET hive.support.concurrency=true;
SET hive.txn.manager=org.apache.hadoop.hive.q1.lockmgr.DbTxnManager;
ALTER TABLE Tab62 SET TBLPROPERTIES ('transactional'='true');
DELETE FROM Tab62 WHERE Tab62.BranchID='Branch2';
```

```
SELECT * FROM Tab62;
```

tab62.branchid	tab62.totalconsumers
Branch1	1115974
Branch3	2591062
Branch4	3940862
Branch5	3941177
Branch6	8032799
Branch7	6566187
Branch8	2583583
Branch9	5208016

Figure: Problem 6

Problem4

 williamgantterry@LAPTOP-H2L556UK: ~

value: ignored

OK

+-----+-----+	
tab62.branchid	tab62.totalconsumers
+-----+-----+	
Branch1	1115974
Branch2	5099141
Branch3	2591062
Branch4	3940862
Branch5	3941177
Branch6	8032799
Branch7	6566187
Branch8	2583583
Branch9	5208016
+-----+-----+	

9 rows selected (0.35 seconds)

Figure: Problem 6

Problem4

 williamgantterry@LAPTOP-H2L556UK: ~

```
0: jdbc:hive2://> 0: jdbc:hive2://>  
SELECT * FROM Tab62;21/08/04 09:54:54 [1890c451  
OK
```

tab62.branchid	tab62.totalconsumers
Branch1	1115974
Branch3	2591062
Branch4	3940862
Branch5	3941177
Branch6	8032799
Branch7	6566187
Branch8	2583583
Branch9	5208016

```
8 rows selected (0.271 seconds)
```

GitHub

<https://github.com/williamgantterry/P1>